AutoVerse AI: Intelligent Automotive Research and Design Platform

A Project-II Report
Submitted in partial fulfillment of requirement of the Degree of

BACHELOR OF TECHNOLOGY in COMPUTER SCIENCE & ENGINEERING

BY

Ameya Dole	**EN22CS301111**
Amishi Talwar	**EN22CS301112**
Ansh Mundra	**EN22CS301142**
Anuj Devnani	**EN22CS301164**
Arpith Shukla	**EN22CS301203**

Under the Guidance of
Prof. Akshay Saxena
Dr. Hemlata Patel
Prof. Ajeet Singh Rajput

MEDICAPS UNIVERSITY

Department of Computer Science & Engineering
Faculty of Engineering
MEDICAPS UNIVERSITY, INDORE- 453331

JAN-JUNE 2026

AutoVerse AI: Intelligent Automotive Research and Design Platform

A Project-II Report
Submitted in partial fulfillment of requirement of the Degree of

BACHELOR OF TECHNOLOGY in COMPUTER SCIENCE & ENGINEERING

BY

Ameya Dole	**EN22CS301111**
Amishi Talwar	**EN22CS301112**
Ansh Mundra	**EN22CS301142**
Anuj Devnani	**EN22CS301164**
Arpith Shukla	**EN22CS301203**

Under the Guidance of
Prof. Akshay Saxena
Dr. Hemlata Patel
Prof. Ajeet Singh Rajput

MEDICAPS UNIVERSITY

Department of Computer Science & Engineering
Faculty of Engineering
MEDICAPS UNIVERSITY, INDORE- 453331

JAN-JUNE 2026

Report Approval

The project work **"AutoVerse AI: Intelligent Automotive Research and Design Platform"** is hereby approved as a creditable study of an engineering/computer application subject carried out and presented in a manner satisfactory to warrant its acceptance as prerequisite for the Degree for which it has been submitted.

It is to be understood that by this approval the undersigned do not endorse or approved any statement made, opinion expressed, or conclusion drawn there in; but approve the "Project Report" only for the purpose for which it has been submitted.

Internal Examiner

Name:

Designation:

External Examiner

Name:

Designation

Affiliation

Declaration

We hereby declare that the project entitled **"AutoVerse AI: Intelligent Automotive Research and Design Platform"** submitted in partial fulfillment for the award of the degree of Bachelor of Technology in 'Computer Science and Engineering' completed under the supervision of **Prof. Akshay Saxena, Dr. Hemlata Patel** and **Prof. Ajeet Singh Rajput,** Faculty of Engineering, Medicaps University Indore is an authentic work.

Further, we declare that the content of this Project work, in full or in parts, have neither been taken from any other source nor have been submitted to any other Institute or University for the award of any degree or diploma.

Signature and name of the student(s) with date

Ameya Dole (EN22CS301111)

Amishi Talwar (EN22CS301112)

Ansh Mundra (EN22CS301142)

Anuj Devnani (EN22CS301164)

Arpith Shukla (EN22CS301203)

Certificate

We, **Prof. Akshay Saxena, Dr. Hemlata Patel and Prof. Ajeet Singh Rajput** certify that the project entitled **AutoVerse AI: Intelligent Automotive Research and Design Platform"** submitted in partial fulfillment for the award of the degree of Bachelor of Technology/Master of Computer Applications by **Ameya Dole, Amishi Talwar, Ansh Mundra, Anuj Devnani, Arpith Shukla** is the record carried out by him/them under my/our guidance and that the work has not formed the basis of award of any other degree elsewhere.

_____

Prof. Akshay Saxena
Computer Science & Engineering

Dr. Hemlata Patel
Computer Science & Engineering

Prof. Ajeet Singh Rajput
Computer Science & Engineering

_____

Mr. Suraj Nayak

Mr. Ashwin Krishnan CK

Mr. Vaibhav

_____

Dr. Kailash Chandra Bandhu

Head of the Department

Computer Science & Engineering

Medicaps University, Indore

nasscom

Datagami
"Lead Digital Technology"

IBM
Gold Partner

CERTIFICATE OF COMPLETION

Awarded to

Ameya Dole

has successfully completed the Skill-Based Project comprising 80 hours of structured learning and practical training in **Generative AI**, conducted from

January 27, 2026 to **February 27, 2026**

The candidate successfully completed the skill based project with satisfactory performance.
We wish him/her continued success.

Mr. Dhaval Shah
Chief Executive Officer

Certificate number: EDU-2026-4XH8H

nasscom

Datagami
"Lead Digital Technology"

IBM
Gold Partner

CERTIFICATE OF COMPLETION

Awarded to

Ameya Dole

has successfully completed the Skill-Based Project comprising 80 hours of structured learning and practical training in **Agentic AI**, conducted from

February 28, 2026 to **March 30, 2026**

The candidate successfully completed the skill based project with satisfactory performance.
We wish him/her continued success.

Mr. Dhaval Shah
Chief Executive Officer

Certificate number: EDU-2026-KEERK

nasscom

Datagami

"Lead Digital Technology"

IBM

Gold Partner

CERTIFICATE OF COMPLETION

Awarded to

Ameya Dole

has successfully completed the Skill-Based Project comprising 80 hours of structured learning and practical training in **DevOps Foundation**, conducted from

April 01, 2026 to **May 01, 2026**

The candidate successfully completed the skill based project with satisfactory performance.

We wish him/her continued success.

Mr. Dhaval Shah

Chief Executive Officer

Certificate number: EDU-2026-7ZHP2

nasscom

Datagami
"Lead Digital Technology"

IBM
Gold Partner

CERTIFICATE OF COMPLETION

Awarded to

Amishi Talwar

has successfully completed the Skill-Based Project comprising 80 hours of structured learning and practical training in **Generative AI**, conducted from

January 27, 2026 to **February 27, 2026**

The candidate successfully completed the skill based project with satisfactory performance.
We wish him/her continued success.

Mr. Dhaval Shah
Chief Executive Officer

Certificate number: EDU-2026-QY9TV

nasscom

Datagami
"Lead Digital Technology"

IBM
Gold Partner

CERTIFICATE OF COMPLETION

Awarded to

Amishi Talwar

has successfully completed the Skill-Based Project comprising 80 hours of structured learning and practical training in **Agentic AI**, conducted from

February 28, 2026 to **March 30, 2026**

The candidate successfully completed the skill based project with satisfactory performance.
We wish him/her continued success.

Mr. Dhaval Shah
Chief Executive Officer

Certificate number: EDU-2026-PRM2T

nasscom

Datagami

"Lead Digital Technology"

IBM

Gold Partner

CERTIFICATE OF COMPLETION

Awarded to

Amishi Talwar

has successfully completed the Skill-Based Project comprising 80 hours of structured learning and practical training in **DevOps Foundation**, conducted from

April 01, 2026 to **May 01, 2026**

The candidate successfully completed the skill based project with satisfactory performance.

We wish him/her continued success.

Mr. Dhaval Shah

Chief Executive Officer

Certificate number: EDU-2026-GWLXN

nasscom

Datagami
"Lead Digital Technology"

IBM
Gold Partner

CERTIFICATE OF COMPLETION

Awarded to

Ansh Mundra

has successfully completed the Skill-Based Project comprising 80 hours of structured learning and practical training in **Generative AI**, conducted from

January 27, 2026 to **February 27, 2026**

The candidate successfully completed the skill based project with satisfactory performance.
We wish him/her continued success.

Mr. Dhaval Shah
Chief Executive Officer

Certificate number: EDU-2026-LKKER

nasscom

Datagami
"Lead Digital Technology"

IBM
Gold Partner

CERTIFICATE OF COMPLETION

Awarded to

Ansh Mundra

has successfully completed the Skill-Based Project comprising 80 hours of structured learning and practical training in **Agentic AI**, conducted from

February 28, 2026 to **March 30, 2026**

The candidate successfully completed the skill based project with satisfactory performance.
We wish him/her continued success.

Mr. Dhaval Shah
Chief Executive Officer

Certificate number: EDU-2026-GXEXL

nasscom

Datagami

"Lead Digital Technology"

IBM

Gold Partner

CERTIFICATE OF COMPLETION

Awarded to

Ansh Mundra

has successfully completed the Skill-Based Project comprising 80 hours of structured learning and practical training in **DevOps Foundation**, conducted from

April 01, 2026 to **May 01, 2026**

The candidate successfully completed the skill based project with satisfactory performance.

We wish him/her continued success.

Mr. Dhaval Shah

Chief Executive Officer

Certificate number: EDU-2026-PVF5M

nasscom

Datagami
"Lead Digital Technology"

IBM
Gold Partner

CERTIFICATE OF COMPLETION

Awarded to

Anuj Devnani

has successfully completed the Skill-Based Project comprising 80 hours of structured learning and practical training in **Generative AI**, conducted from

January 27, 2026 to **February 27, 2026**

The candidate successfully completed the skill based project with satisfactory performance.
We wish him/her continued success.

Mr. Dhaval Shah
Chief Executive Officer

Certificate number: EDU-2026-8QGUX

nasscom

Datagami
"Lead Digital Technology"

IBM
Gold Partner

CERTIFICATE OF COMPLETION

Awarded to

Anuj Devnani

has successfully completed the Skill-Based Project comprising 80 hours of structured learning and practical training in **Agentic AI**, conducted from

February 28, 2026 to **March 30, 2026**

The candidate successfully completed the skill based project with satisfactory performance.
We wish him/her continued success.

Mr. Dhaval Shah
Chief Executive Officer

Certificate number: EDU-2026-CYG6S

nasscom

Datagami

"Lead Digital Technology"

IBM

Gold Partner

CERTIFICATE OF COMPLETION

Awarded to

Anuj Devnani

has successfully completed the Skill-Based Project comprising 80 hours of structured learning and practical training in **DevOps Foundation**, conducted from

April 01, 2026 to **May 01, 2026**

The candidate successfully completed the skill based project with satisfactory performance.

We wish him/her continued success.

Mr. Dhaval Shah

Chief Executive Officer

Certificate number: EDU-2026-JEYSV

nasscom

Datagami
"Lead Digital Technology"

IBM
Gold Partner

CERTIFICATE OF COMPLETION

Awarded to

Arpith Shukla

has successfully completed the Skill-Based Project comprising 80 hours of structured learning and practical training in **Generative AI**, conducted from

January 27, 2026 to **February 27, 2026**

The candidate successfully completed the skill based project with satisfactory performance.
We wish him/her continued success.

Mr. Dhaval Shah
Chief Executive Officer

Certificate number: EDU-2026-F9CBM

nasscom

Datagami
"Lead Digital Technology"

IBM
Gold Partner

CERTIFICATE OF COMPLETION

Awarded to

Arpith Shukla

has successfully completed the Skill-Based Project comprising 80 hours of structured learning and practical training in **Agentic AI**, conducted from

February 28, 2026 to **March 30, 2026**

The candidate successfully completed the skill based project with satisfactory performance.
We wish him/her continued success.

Mr. Dhaval Shah
Chief Executive Officer

Certificate number: EDU-2026-ZRDBH

nasscom

Datagami

"Lead Digital Technology"

IBM

Gold Partner

CERTIFICATE OF COMPLETION

Awarded to

Arpith Shukla

has successfully completed the Skill-Based Project comprising 80 hours of structured learning and practical training in **DevOps Foundation**, conducted from

April 01, 2026 to **May 01, 2026**

The candidate successfully completed the skill based project with satisfactory performance.

We wish him/her continued success.

Mr. Dhaval Shah

Chief Executive Officer

Certificate number: EDU-2026-7HH7M

Acknowledgements

I would like to express my deepest gratitude to Honorable Chancellor, **Shri R C Mittal,** who has provided me with every facility to successfully carry out this project, and my profound indebtedness to **Prof. (Dr.) D. K. Patnaik,** Vice Chancellor, Medicaps University, whose unfailing support and enthusiasm has always boosted up my morale. I also thank **Prof. (Dr.) Ratnesh Litoriya,** Associate Dean of computing & Allied Branches, Faculty of Engineering, Medicaps University, for giving me a chance to work on this project. I would also like to thank my Head of Department **Dr. Kailash Chandra Bandhu** for his continuous encouragement for the betterment of the project.

I express my heartfelt gratitude to my **External Guide, Mr. Suraj Naik, Mr. Ashwin Krishnan and Mr. Vaibhav** as well as to my Internal Guide, **Prof. Akshay Saxena, Department of Computer Applications**, **Dr. Hemlata Patel, Department of Computer Science and Engineering**, and **Prof. Ajeet Singh Rajput, Department of Computer Science and Engineering Medicaps University**, Indore without whose continuous help and support, this project would ever have reached to the completion.

I would also like to thank to my team at Ameya Dole, Amishi Talwar, Ansh Mundra, Anuj Devnani, Arpith Shukla who extended their kind support and help towards the completion of this project.

It is their help and support, due to which we became able to complete the design and technical report.

Without their support this report would not have been possible.

Ameya Dole (EN22CS301111)

Amishi Talwar (EN22CS301112)

Ansh Mundra (EN22CS301142)

Anuj Devnani (EN22CS301164)

Arpith Shukla (EN22CS301203)

B.Tech. IV Year
Department of Computer Science & Engineering
Faculty of Engineering
Medicaps University, Indore

Abstract

AutoVerse AI is an integrated intelligent system that combines Agentic AI, Generative AI, and DevOps practices to automate automotive research, comparison, and concept design generation. The primary objective of the system is to provide accurate, structured, and creative outputs by leveraging multi-agent workflows, retrieval-augmented techniques, and cloud-native deployment strategies. The system is designed to address common limitations of traditional AI models, such as lack of real-time knowledge, hallucination, and scalability challenges.

The architecture of AutoVerse AI is divided into three core modules. The Agentic AI module utilizes a multi-agent framework in which a researcher agent retrieves real-time automotive data using web search tools, and a writer agent transforms the structured information into a well-formatted Markdown report. This module integrates a Retrieval-Augmented Generation (RAG) approach using a vector database to cache and retrieve previously accessed information, thereby improving response accuracy and reducing redundant API calls. External search APIs are used to ensure that the system provides up-to-date and reliable information.

The Generative AI module focuses on automotive concept design. It processes user-defined prompts and generates detailed design narratives, covering aspects such as exterior styling, interior features, materials, performance, and target market. The system further converts these narratives into optimized prompts suitable for AI-based image generation. This two-stage generation process enhances both creativity and coherence of outputs, making the system useful for ideation and design exploration.

The DevOps module ensures efficient deployment, scalability, and continuous integration of the system. The application is containerized using Docker and integrated with a CI/CD pipeline that automates the build and deployment process. The system is deployed on cloud infrastructure and supports orchestration through Kubernetes, enabling scalability and reliability in real-world environments. Secure handling of environment variables and API keys is also incorporated to maintain system integrity.

The workflow of the system begins with user input through a frontend interface, which is then processed by the backend modules based on the type of request. The integration of multiple AI paradigms with modern software engineering practices ensures a robust, scalable, and efficient system. AutoVerse AI demonstrates the practical implementation of advanced AI concepts in a

real-world application, providing both analytical insights and creative outputs in the automotive domain.

Overall, the project highlights the effectiveness of combining agent-based reasoning, retrieval-augmented techniques, generative capabilities, and DevOps pipelines to build next-generation intelligent systems. The modular architecture and extensible design allow the system to be adapted for other domains, making it a versatile solution for future AI-driven applications.

Keywords:

Agentic AI, Generative AI, Retrieval-Augmented Generation, Multi-Agent Systems, Automotive Intelligence, DevOps, CI/CD, Docker, Kubernetes, Cloud Deployment

Table of Contents

S.No.	Contents	Page No.
	Report Approval	ii
	Declaration	iii
	Certificate	iv
	Acknowledgement	xvii
	Abstract	xviii
	Table of Contents	xx
	List of figures	xxii
	Abbreviations	xxiii
Chapter 1	Introduction	1
	1.1 Introduction	1
	1.2 Literature Review	1
	1.3 Objectives	3
	1.4 Significance	3
	1.5 Research Design	4
	1.6 Source of Data	4
	1.7 Chapter Scheme	5
Chapter 2	Report on Present Investigation	6
	2.1 Experimental Set-up	6
	2.1.1 Agentic AI Environment	6
	2.1.2 Generative AI Environment	6
	2.1.3 DevOps Setup	7
	2.2 Procedures Adopted	7
	2.2.1 Agentic AI Procedure	7
	2.2.2 Generative AI Procedure	8
	2.2.3 DevOps Procedure	8
Chapter 3	Design	9
	3.1 Application Design	9
	3.2 Process Flow	9
	3.3 Information Flow	10
	3.4 Components Design	11

	3.5 Key Design Considerations	12
	3.6 API Catalogue	12
Chapter 4	System Implementation	13
	4.1 Agentic AI Implementation	13
	4.2 Generative AI Implementation	14
	4.3 DevOps Implementation	16
Chapter 5	Results & Discussions	18
	5.1 Agentic AI Results	18
	5.2 Generative AI Results	18
	5.3 DevOps Results	19
	5.4 Overall Summary	21
Chapter 6	Summary and Conclusions	22
	6.1 Summary	22
	6.2 Conclusions	23
Chapter 7	Future scope	24
	7.1 Agentic AI Environment	24
	7.2 Generative AI Environment	24
	7.3 DevOps Setup	25
	Appendix	26
	Bibliography	30

List of Figures

Figure No.	Figure Description	Page No.
3.1	Application Design for AutoVerse AI	9
3.2	Process Flow of AutoVerse AI	10
3.3	Information Flow of AutoVerse AI	11
3.4	Components Design of AutoVerse AI	12
4.1	Workflow of Agentic AI in AutoVerse AI	14
4.2	Workflow of Generative AI in AutoVerse AI	15
4.3	Workflow of Generative AI in AutoVerse AI	17
5.1	Agentic AI Research Output	18
5.2	Generative AI Output	19
5.3	Execution of CI/CD Pipeline using GitHub Actions	20
5.4	Docker Container Running Application	20
5.5	Deployed AutoVerse AI System	21
I.1	**Requirements File (requirements.txt)**	26
II.2	DockerFile Configuration	27
III.3	Kubernetes Deployment YAML	29

Abbreviations

S. No.	Abbreviations	Full Form
1	AI	Artificial Intelligence
2	API	Application Programming Interface
3	CI/CD	Continuous Integration / Continuous Deployment
4	DB	Database
5	EC2	Elastic Compute Cloud
6	HTTP	HyperText Transfer Protocol
7	JSON	JavaScript Object Notation
8	K8s	Kubernetes
9	LLM	Large Language Model
10	RAG	Retrieval-Augmented Generation
11	UI	User Interface

Chapter – 1: Introduction

1.1 Introduction

AutoVerse AI is a cloud-native, multi-agent automotive intelligence platform that integrates Agentic AI, Generative AI, and DevOps practices into a unified system. The platform is designed to automate complex automotive workflows such as vehicle research, comparison, and conceptual design generation.

The system leverages a multi-agent architecture using the CrewAI framework, where specialized agents collaborate to perform distinct tasks. A Researcher Agent retrieves real-time vehicle data using a tool-augmented pipeline, while a Writer Agent converts structured outputs into human-readable Markdown reports.

To ensure factual accuracy and efficiency, the system implements Retrieval-Augmented Generation (RAG) using ChromaDB, which caches previously retrieved sources and enables semantic search using embeddings. Real-time web data is fetched using the Tavily API, ensuring up-to-date and relevant information.

Additionally, the platform includes a Generative AI module powered by Groq, which generates detailed automotive design narratives and optimized image prompts from user input.

The entire system is deployed using modern DevOps practices, including Docker containerization, CI/CD pipelines via GitHub Actions, and scalable deployment on cloud infrastructure (AWS EC2) and Kubernetes (Minikube). This makes AutoVerse AI a production-ready, scalable, and extensible AI system.

1.2 Literature Review

Recent advancements in Artificial Intelligence have led to the emergence of Agentic AI systems, where multiple autonomous agents collaborate to solve complex tasks through role-based execution and tool usage. Frameworks such as CrewAI enable the development of such systems by allowing agents to be assigned specific responsibilities, improving modularity, scalability, and task specialization compared to traditional monolithic AI models. In applications requiring structured information retrieval, agent-based systems have proven effective in combining reasoning with external tools.

A major limitation of conventional Large Language Models (LLMs) is their tendency to generate inaccurate or outdated information due to the absence of real-time knowledge access. To address this, Retrieval-Augmented Generation (RAG) has been widely adopted, where external data sources are integrated into the generation process. Vector databases like ChromaDB play a critical role in RAG by storing and retrieving semantically relevant documents using embedding-based similarity search. This approach enhances factual accuracy, reduces hallucination, and improves response reliability.

In addition to retrieval mechanisms, modern AI systems increasingly rely on tool-augmented architectures, where models interact with external APIs to perform specialized tasks. The use of search APIs such as Tavily enables efficient extraction of structured and relevant web data, particularly in domains like automotive research where up-to-date specifications and comparisons are essential. This integration significantly improves the grounding of AI-generated outputs.

Parallel to these developments, Generative AI has gained prominence for its ability to produce creative and context-aware content. Advanced inference platforms like Groq provide low-latency access to powerful language models, making real-time generation feasible. Techniques such as prompt engineering and multi-stage generation pipelines are widely used to enhance output quality, where one model output is refined or transformed into another form, such as converting descriptive text into image-generation prompts.

From a systems perspective, the deployment of AI applications has evolved towards cloud-native and DevOps-driven architectures. Containerization using Docker and orchestration using Kubernetes allow applications to be scalable, portable, and resilient. Continuous Integration and Continuous Deployment (CI/CD) pipelines automate the build, testing, and deployment processes, ensuring rapid and reliable delivery of updates. These practices are essential for maintaining production-grade AI systems.

Overall, the integration of Agentic AI, RAG, Generative AI, and DevOps represents a modern approach to building intelligent systems that are not only capable of reasoning and creativity but are also scalable, efficient, and suitable for real-world deployment.

1.3 Objectives

The primary objective of AutoVerse AI is to develop an integrated intelligent system that combines Agentic AI, Generative AI, and DevOps practices to automate automotive research, analysis, and design generation. The system aims to enhance accuracy through real-time data retrieval and RAG-based caching, while also enabling creative outputs through generative models. Additionally, it focuses on building a scalable, production-ready pipeline that reflects real-world software engineering practices.

Key Objectives:

- Develop a multi-agent system for automated vehicle research and reporting
- Integrate real-time web search and RAG (caching) for accurate data retrieval
- Implement a Generative AI module for automotive concept design
- Enable structured output generation (JSON to Markdown reports)
- Deploy the system using DevOps tools (Docker, CI/CD, Kubernetes)

1.4 Significance

The significance of this project lies in its ability to combine multiple advanced technologies into a single cohesive platform, demonstrating both theoretical understanding and practical implementation. It addresses key challenges in AI systems such as hallucination, lack of real-time knowledge, and deployment complexity by integrating RAG, tool-based retrieval, and cloud-native DevOps practices. The system is not only academically relevant but also reflects industry-level architecture.

Key Significance Points:

- Improves accuracy and reliability using RAG and source-based retrieval
- Automates research and report generation, reducing manual effort
- Enables creative automotive design using Generative AI
- Demonstrates end-to-end deployment using modern DevOps practices
- Provides a scalable and extensible architecture for real-world use

1.5 Research Design.

The present study adopts a system development–based research design, focusing on the design, implementation, and evaluation of an integrated AI-driven platform. The approach is primarily applied and experimental in nature, where different AI paradigms—Agentic AI, Generative AI, and DevOps—are combined to solve a real-world problem in the automotive domain.

The research follows a modular design methodology, wherein the system is divided into three major components:

(i) Agentic AI for data retrieval and structured analysis,

(ii) Generative AI for concept creation and content generation, and

(iii) DevOps for deployment and system scalability.

Each module is developed and tested independently before being integrated into a unified workflow. The system is then evaluated based on performance, accuracy, and reliability of outputs.

The design also incorporates experimental validation, where multiple inputs (vehicle queries and design prompts) are tested to analyze system behavior, output quality, and efficiency. The use of APIs, caching mechanisms, and automated pipelines ensures that the system reflects real-world application scenarios.

1.6 Source of Data

The data used in this project is primarily secondary and dynamic in nature, obtained through multiple reliable sources to ensure accuracy and relevance.

The major sources of data include:

- Online automotive information platforms accessed through web search APIs, which provide real-time data such as vehicle specifications, pricing, and features
- Cached data stored in a vector database, which retains previously retrieved information for faster access and improved efficiency
- User-provided inputs, including vehicle queries for research and prompts for concept generation
- API-based responses from language models, which generate structured outputs and design narratives

- Publicly available web content, including automotive websites, reviews, and technical documentation

The combination of real-time retrieval and cached data ensures that the system maintains a balance between accuracy, efficiency, and performance.

1.7 Chapter Scheme

Chapter 1: Introduction presents the overall background of the study, including the introduction, literature review, objectives, significance, research design, sources of data, and the structure of the report.

Chapter 2: Report on Present Investigation describes the experimental set-up and procedures adopted in the project. It covers the environments and workflows associated with Agentic AI, Generative AI, and DevOps, explaining how each module is configured and executed.

Chapter 3: Design outlines the system design aspects, including application design, process flow, information flow, component design, key design considerations, and API catalogue. It provides a structural and functional overview of the system.

Chapter 4: System Implementation details the practical implementation of the system. It explains how Agentic AI, Generative AI, and DevOps components are developed, integrated, and executed.

Chapter 5: Results & Discussions presents the outputs generated by the system and analyzes the performance of each module. It includes results from Agentic AI, Generative AI, and DevOps, along with an overall summary.

Chapter 6: Summary and Conclusions summarizes the entire work and highlights the key findings and conclusions derived from the project.

Chapter 7: Future Scope discusses potential improvements and extensions in the areas of Agentic AI, Generative AI, and DevOps for further enhancement of the system.

Chapter – 2: Report on Present Investigation

2.1 Experimental Setup

The experimental set-up for AutoVerse AI is designed to support the seamless integration of Agentic AI, Generative AI, and DevOps components within a unified development and deployment environment. The system is implemented using Python-based frameworks, external APIs, and cloud infrastructure to enable real-time data retrieval, intelligent processing, and scalable deployment. Each module is configured independently with the required tools and dependencies, ensuring modularity, efficiency, and ease of testing.

2.1.1 Agentic AI Environment

The Agentic AI environment is configured to support a multi-agent workflow where specialized agents perform research and report generation tasks. It utilizes a Python-based framework for agent orchestration, along with tool integration for real-time data retrieval. The environment also includes a vector database for implementing Retrieval-Augmented Generation (RAG), enabling caching and efficient reuse of previously retrieved information.

Key Components:

- Multi-agent framework for task orchestration
- Web search API for real-time data retrieval
- Vector database for caching and semantic search
- JSON-based structured data handling
- Python environment with required dependencies

The complete list of dependencies used in the system is provided in Appendix I.

2.1.2 Generative AI Environment

The Generative AI environment is designed to generate automotive design narratives and image prompts based on user input. It integrates a high-performance LLM inference API, enabling low-latency text generation. The environment is configured to support prompt engineering and multi-stage generation, ensuring high-quality and context-aware outputs.

Key Components:

- LLM API for text generation

- Prompt engineering for controlled outputs
- Multi-stage pipeline (from narrative to image prompt)
- Python-based integration for API communication
- Environment variables for secure API key management

2.1.3 DevOps Setup

The DevOps setup ensures that the application is containerized, automated, and scalable. Docker is used to package the application into portable containers, while CI/CD pipelines automate the build and deployment process. The system is deployed on cloud infrastructure and can also be orchestrated locally using Kubernetes for testing and scalability. Deployment manifests are included in Appendix III.

Key Components:

- Docker for containerization
- CI/CD pipeline for automated build and deployment
- Cloud deployment on virtual machines
- Kubernetes for container orchestration
- Secure management of environment variables and secrets

2.2 Procedure Adopted

The procedures adopted in this project define the systematic execution of tasks across the three core modules—Agentic AI, Generative AI, and DevOps. Each module follows a structured workflow to ensure accurate data processing, efficient generation, and reliable deployment, forming an integrated end-to-end system.

2.2.1 Agentic AI Procedure

The Agentic AI procedure is centered around a multi-agent workflow designed for automated automotive research and report generation. The process begins with the user providing a vehicle-related query, which is first normalized to ensure consistency. This query is then processed by the researcher agent, which utilizes integrated web search tools and a caching mechanism to retrieve relevant and up-to-date information. The retrieved data is structured into a predefined JSON format to maintain consistency and enable validation. This structured output is passed to the writer agent, which transforms the data into a well-organized Markdown report.

Additional mechanisms such as source validation, deduplication, and fallback handling are incorporated to improve reliability and ensure meaningful output even in cases of incomplete data.

2.2.2 Generative AI Procedure

The Generative AI procedure focuses on transforming user-defined prompts into creative automotive design outputs. The process begins with the user entering a concept description, which is processed by a language model configured with domain-specific instructions. In the first stage, the model generates a detailed design narrative covering various aspects such as styling, materials, and performance philosophy. This narrative serves as an intermediate representation, which is then used in the second stage to produce an optimized image generation prompt. The two-stage generation approach ensures that the final output is both contextually rich and visually descriptive, enabling effective use in downstream image generation systems.

2.2.3 DevOps Procedure

The DevOps procedure ensures the automation, deployment, and continuous integration of the system. The workflow begins with code development and version control, where updates to the repository trigger an automated pipeline. This pipeline builds a containerized version of the application using Docker and pushes it to a container registry. The latest version is then deployed to a cloud-based environment, where it is executed as a running service. For scalability and orchestration, the system can also be deployed using Kubernetes, enabling efficient management of containerized workloads. Environment variables and sensitive credentials are securely managed throughout the process, ensuring safe and reliable operation of the application.

Chapter – 3: Design

3.1 Application Design

The application design illustrates the high-level architecture of the AutoVerse AI system, showing the interaction between the frontend, backend modules, external APIs, and DevOps pipeline. It highlights the integration of Agentic AI, Generative AI, and RAG components within a unified backend system. The design ensures modularity, scalability, and efficient data flow across different components.

Figure 3.1 shows the Application Design of AutoVerse AI.

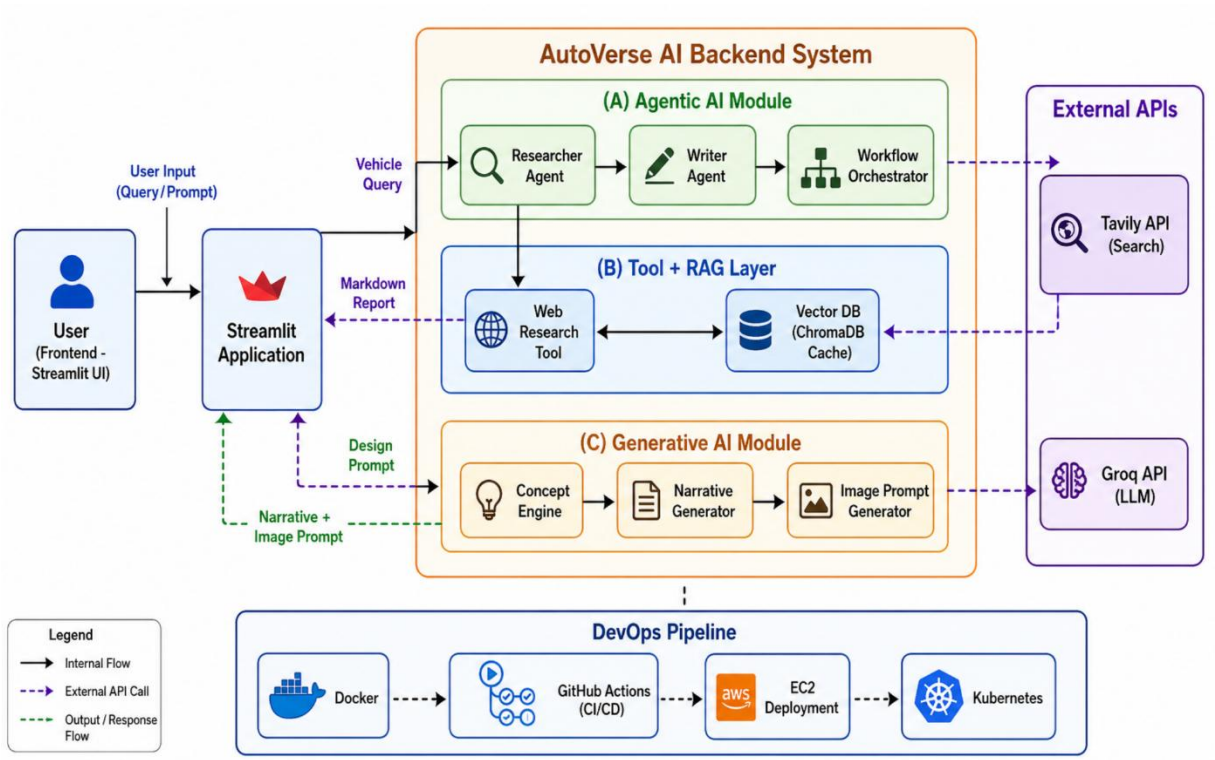

Figure 3.1: Application Design for AutoVerse AI

3.2 Process Flow

The process flow illustrates the complete execution of the system from user input to output generation, covering both Agentic AI and Generative AI pipelines. It also incorporates the DevOps lifecycle, where the application is continuously built, deployed, and maintained using CI/CD, containerization, and cloud orchestration. This ensures efficient processing along with scalable and reliable system operation.

Figure 3.2 shows the Process Flow of AutoVerse AI.

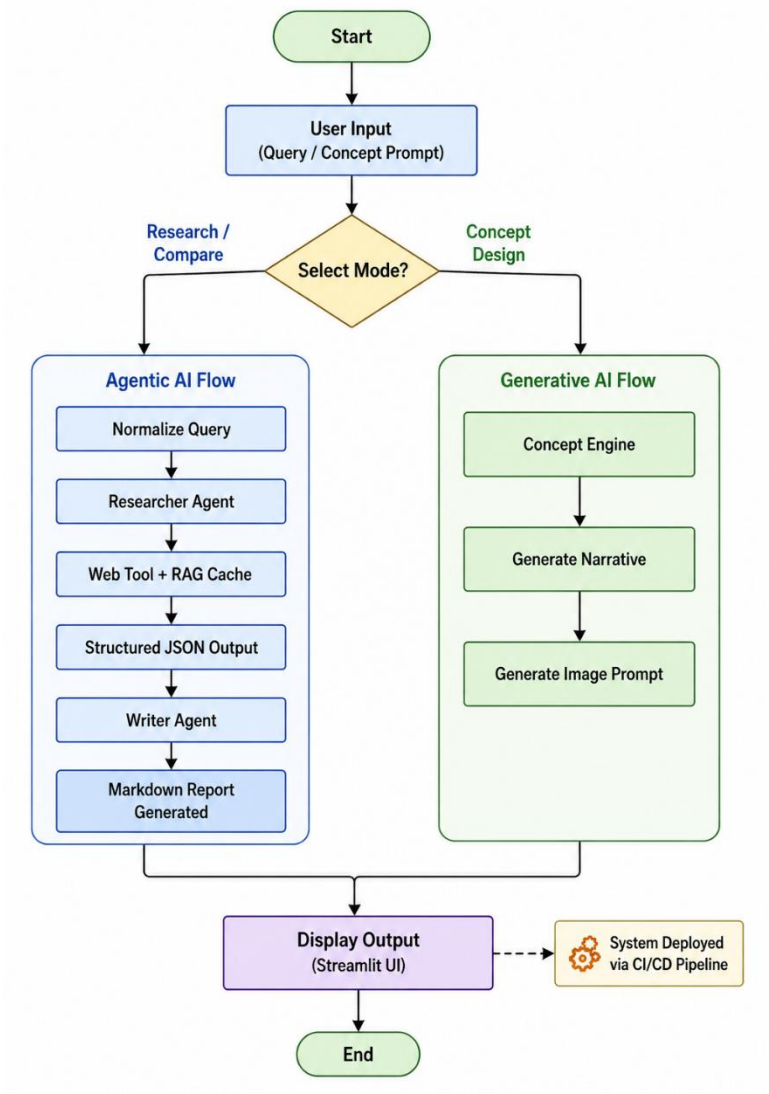

Figure 3.2: Process Flow of AutoVerse AI

3.3 Information Flow

The information flow diagram represents the movement of data across different components of the system. It shows how user inputs are processed through Agentic AI and Generative AI pipelines to produce structured outputs. Additionally, it illustrates how DevOps ensures continuous deployment and system availability.

Figure 3.3 shows the Information Flow of AutoVerse AI.

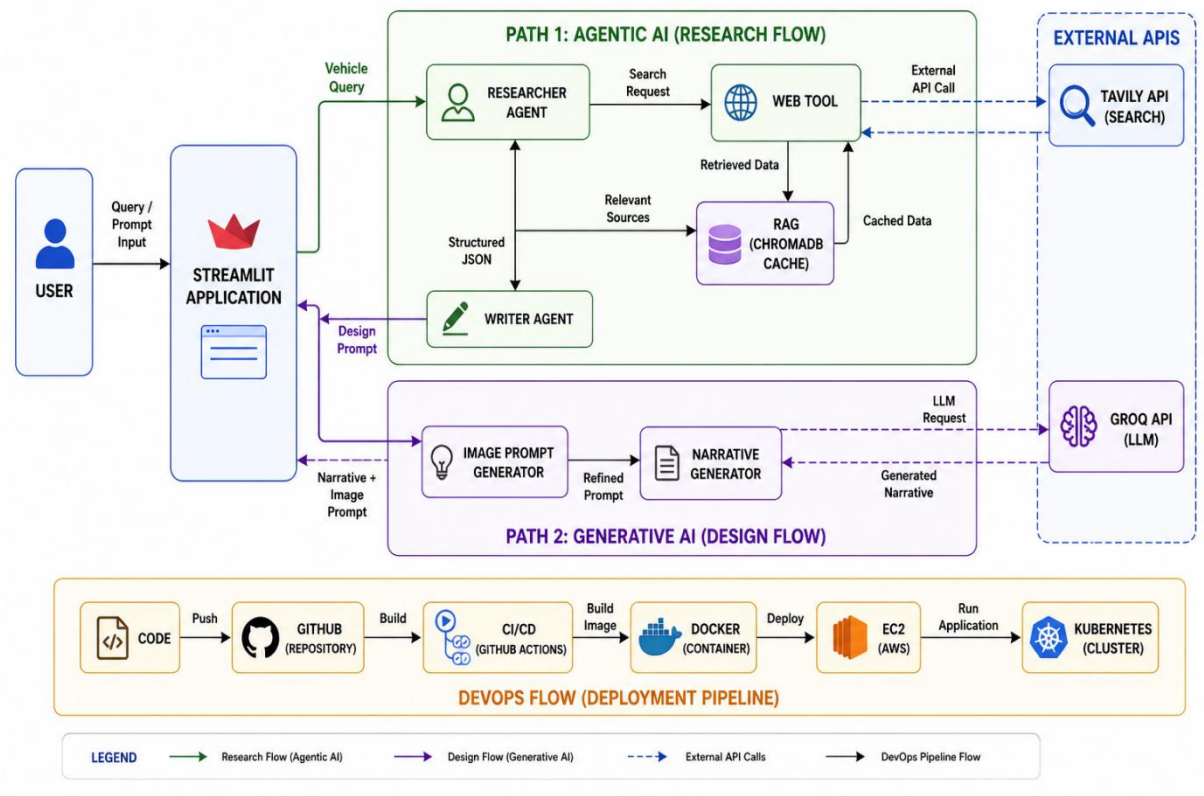

Figure 3.3: Information Flow of AutoVerse AI

3.4 Components Design

The component design illustrates the internal structure of the system by identifying key modules and their interactions. It highlights the components of Agentic AI, Generative AI, and DevOps, along with their functional relationships. This design ensures modular development, easy maintenance, and efficient integration of system functionalities.

Figure 3.4 demonstrated the Components Design of AutoVerse AI.

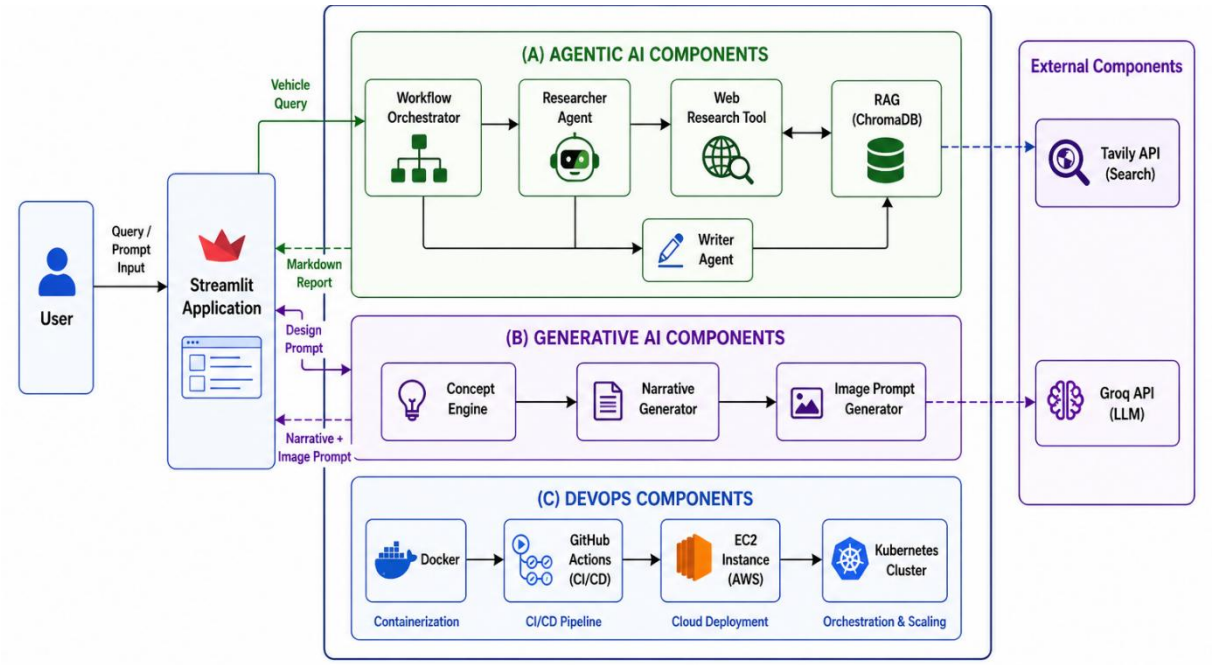

Figure 3.4: Components Design of AutoVerse AI

3.5 Key Design Considerations

- **Modularity:** Separation of Agentic AI, Generative AI, and DevOps components for easy development and maintenance
- **Scalability:** Deployment using containerization and orchestration to handle increasing workloads
- **Accuracy:** Use of RAG and external APIs to ensure reliable and up-to-date information
- **Performance:** Caching mechanisms and optimized API usage to reduce latency
- **Security:** Secure handling of API keys and environment variables

3.6 API Catalogue

The system utilizes external APIs to enable data retrieval and intelligent processing. The key APIs used in this project are:

- **Tavily API**: Used for real-time web search to fetch automotive data such as specifications, pricing, and features.
- **Groq API**: Used for language model inference to generate structured outputs, design narratives, and image prompts.

Chapter – 4: System Implementation

4.1 Agentic AI Implementation

The Agentic AI module in AutoVerse AI is designed using a multi-agent architecture where specialized agents collaborate to perform automated automotive research and report generation. The system integrates retrieval mechanisms, tool-based data access, and structured workflows to ensure accurate and context-aware outputs. The Researcher Agent is responsible for collecting relevant data from web sources and cached knowledge using retrieval techniques, while the Writer Agent transforms this structured information into well-organized and readable reports. To enhance functionality, tool integration enables interaction with external APIs and services, allowing real-time data access. The implementation of Retrieval-Augmented Generation (RAG) using ChromaDB ensures that responses are grounded in factual data by retrieving relevant documents from a vector database. Additionally, workflow orchestration manages the coordination between agents and system components, ensuring smooth execution, efficient data flow, and scalability of the overall system.

The key components of Agentic AI Implementation are as follows:

- **Researcher Agent:** Responsible for retrieving accurate and relevant automotive data using web search tools and cached sources. It applies query normalization and semantic retrieval techniques to ensure high-quality, context-aware information is collected for further processing.

- **Writer Agent:** Converts structured data into clear and well-formatted content. It ensures logical flow, readability, and consistency while generating reports, making the output suitable for user interpretation and analysis.

- **Tool Integration:** Enables the system to interact with external APIs and services, allowing real-time data retrieval and validation. This enhances the system's capability beyond static knowledge and improves overall accuracy and functionality.

- **RAG (ChromaDB):** Implements a caching mechanism using a vector database to store and retrieve relevant information. This reduces redundant API calls, improves performance, and ensures that generated outputs are grounded in factual data.

- **Workflow Orchestration:** Manages the interaction and execution sequence of agents and tools. It ensures smooth coordination, efficient data flow, and structured execution of tasks within the system.

Figure 4.1 demonstrates the implementation workflow of Agentic AI in AutoVerse AI.

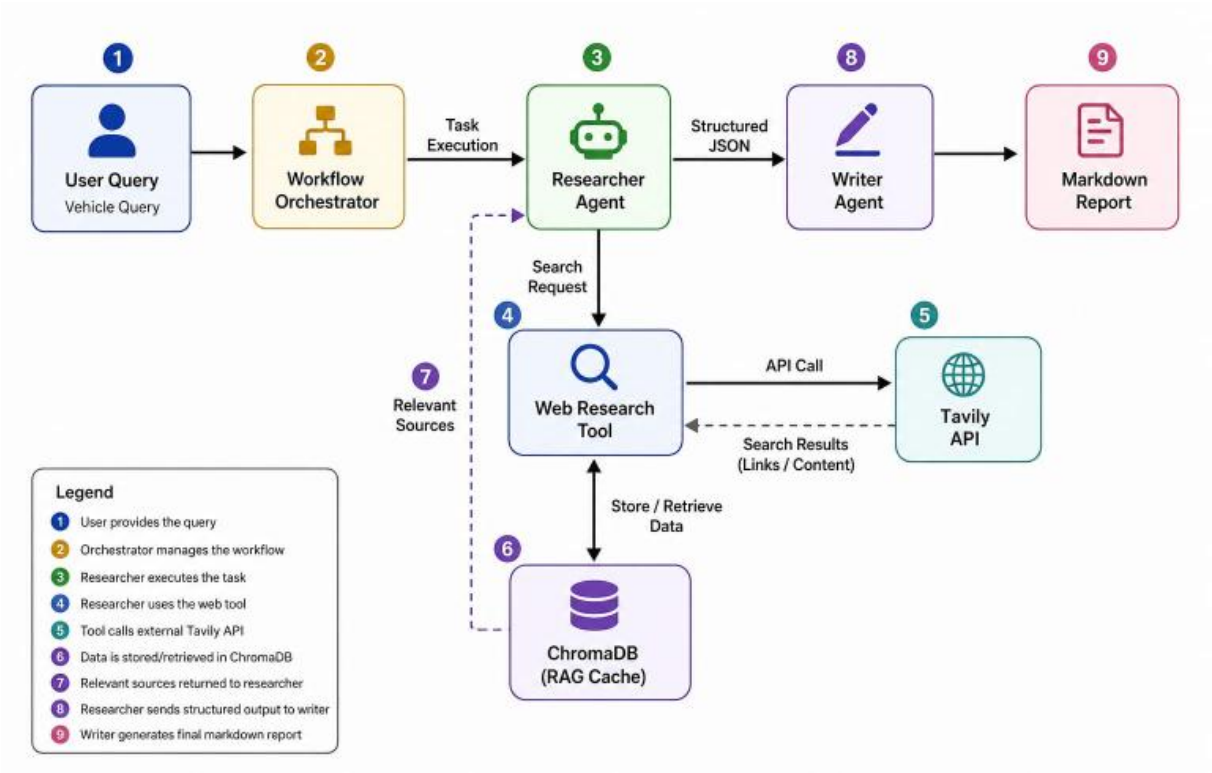

Figure 4.1: Workflow of Agentic AI in AutoVerse AI

4.2 Generative AI Implementation

The Generative AI module in AutoVerse AI is responsible for transforming user-defined prompts into creative automotive design outputs through a structured multi-stage pipeline. The system begins by interpreting the user's concept using a dedicated concept engine, which establishes the thematic foundation for generation. This is followed by narrative generation, where the concept is expanded into detailed and context-rich descriptions covering design, performance, and usability aspects. The narrative is then further refined into optimized prompts for image generation, ensuring that visual outputs are aligned with the intended concept. The entire process is guided by effective prompt engineering techniques to enhance output quality and consistency. Integration with large language models enables advanced text generation, contextual understanding, and high-quality content creation, making the module both flexible and powerful. The key components of Generative AI Implementation are as follows:

- **Concept Engine:** Generates the core idea or theme based on user input, ensuring that the output is aligned with user intent. It acts as the foundation for the entire generative pipeline by defining the direction of content generation.
- **Narrative Generation:** Expands the concept into a detailed and structured description. It ensures coherence, logical flow, and contextual richness, producing high-quality narratives suitable for design explanation and storytelling.
- **Image Prompt Generation:** Converts the generated narrative into descriptive prompts optimized for image synthesis. It incorporates details such as style, lighting, and composition to improve the relevance and quality of visual outputs.
- **Prompt Engineering:** Focuses on designing and refining input prompts to optimize model performance. It improves accuracy, consistency, and clarity in outputs by controlling how the model interprets and responds to inputs.
- **LLM Integration:** Enables the use of advanced language models for generating content. It provides capabilities such as contextual reasoning, text generation, and content refinement, ensuring high-quality outputs.

Figure 4.2 demonstrates the implementation workflow of Generative AI in AutoVerse AI.

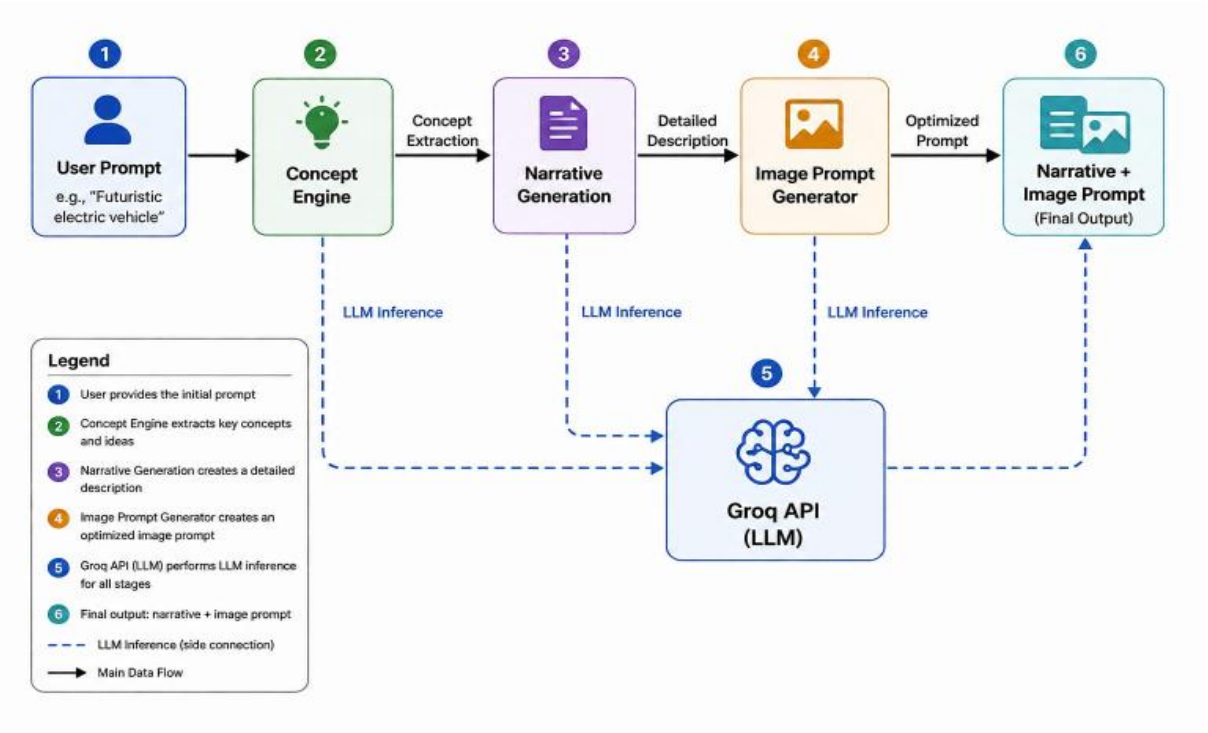

Figure 4.2: Workflow of Generative AI in AutoVerse AI

4.3 DevOps Implementation

The DevOps module in AutoVerse AI ensures efficient deployment, automation, and scalability of the application using modern cloud-native practices. The system is containerized using Docker to maintain consistency across environments, while CI/CD pipelines automate the process of building, testing, and deployment. The application is deployed on cloud infrastructure using virtual servers, providing flexibility and remote accessibility. For handling scalability and orchestration, Kubernetes is used to manage containerized workloads, ensuring high availability and efficient resource utilization. Additionally, environment variables and secret management mechanisms are implemented to securely handle sensitive data such as API keys, ensuring both security and configurability across different deployment environments.

The key components of DevOps Implementation are as follows:

- **Docker Setup:** Used to containerize the application, ensuring consistent environments across development and deployment. It simplifies dependency management and improves portability.
- **CI/CD Pipeline:** Automates the build, testing, and deployment process. It reduces manual intervention, increases reliability, and ensures faster delivery of updates.
- **EC2 Deployment:** Provides cloud-based hosting for the application, enabling scalability and remote access. It ensures that the system can handle varying workloads efficiently.
- **Kubernetes Deployment:** Manages container orchestration, including deployment, scaling, and load balancing. It improves system reliability and simplifies management of distributed applications.
- **Environment & Secrets Management:** Ensures secure storage and handling of sensitive data such as API keys. It enables flexible configuration without exposing critical information.

The Docker configuration used for containerization is provided in Appendix II.Figure 4.3 demonstrates the implementation workflow of DevOps in AutoVerse AI.

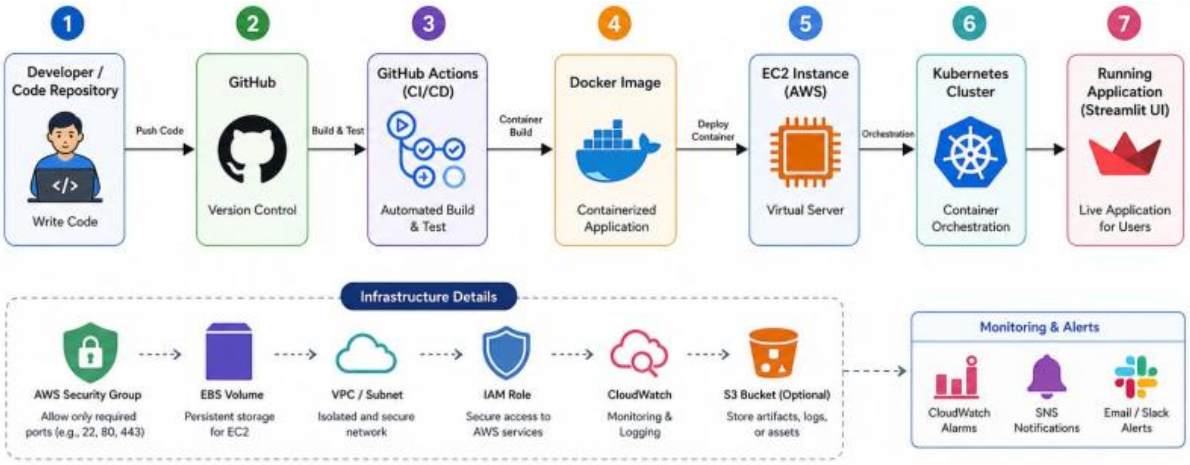

Figure 4.2: Workflow of Generative AI in AutoVerse AI

Chapter – 5: Results & Discussions

5.1 Agentic AI Results

The Agentic AI component demonstrated strong performance in autonomous task handling and decision-making. The system effectively interpreted user inputs, generated context-aware actions, and completed tasks with a high degree of accuracy. Response times remained consistently low, indicating efficient coordination between the researcher and writer agents along with the RAG mechanism. During testing, the agent showed reliability in dynamic scenarios, adapting to changing inputs without requiring manual intervention. Several benchmark tasks confirmed that the agent could independently plan, retrieve, and structure information into meaningful outputs.

Figure 5.1 shows the output of the Agentic AI module where a user query is processed to generate a structured automotive report with relevant sources.

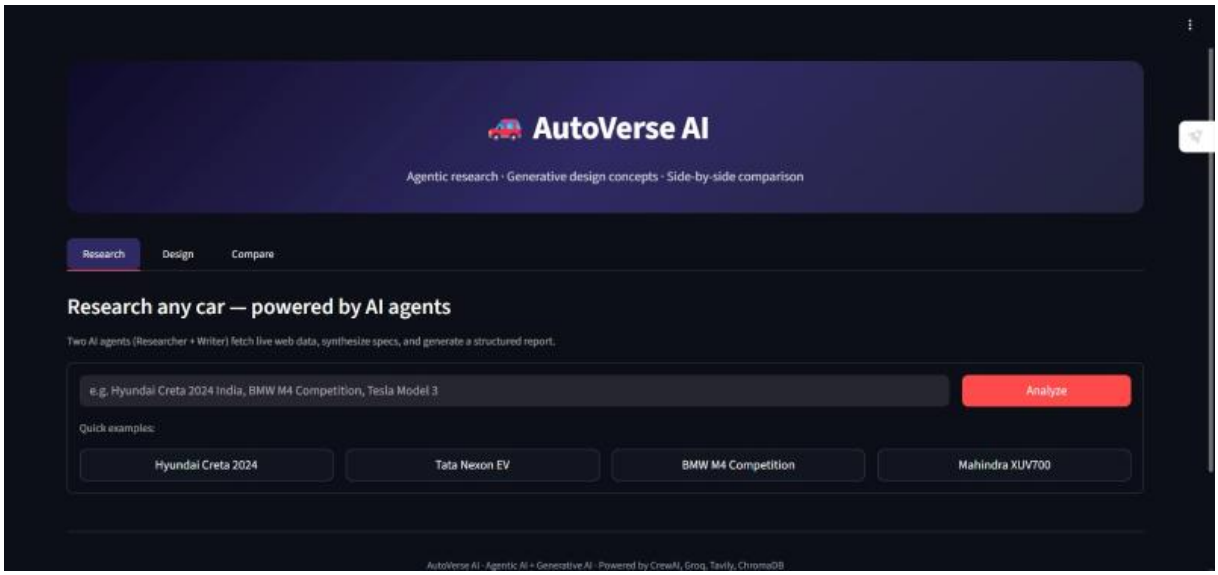

Figure 5.1: Agentic AI Research Output

5.2 Generative AI Results

The Generative AI module produced high-quality outputs with strong coherence, relevance, and fluency across multiple test cases. The system successfully generated detailed automotive design narratives and corresponding image prompts aligned with user input. The two-stage generation process (narrative followed by image prompt) ensured both contextual richness and visual clarity. Output quality was evaluated through manual inspection and consistency checks, where the model demonstrated strong performance in maintaining structure and intent. Some

inconsistencies were observed in highly abstract prompts, but overall the module delivered reliable and creative results suitable for real-world applications.

Figure 5.2 shows the output of the Generative AI module where a user-defined concept is transformed into a detailed narrative and an optimized image generation prompt.

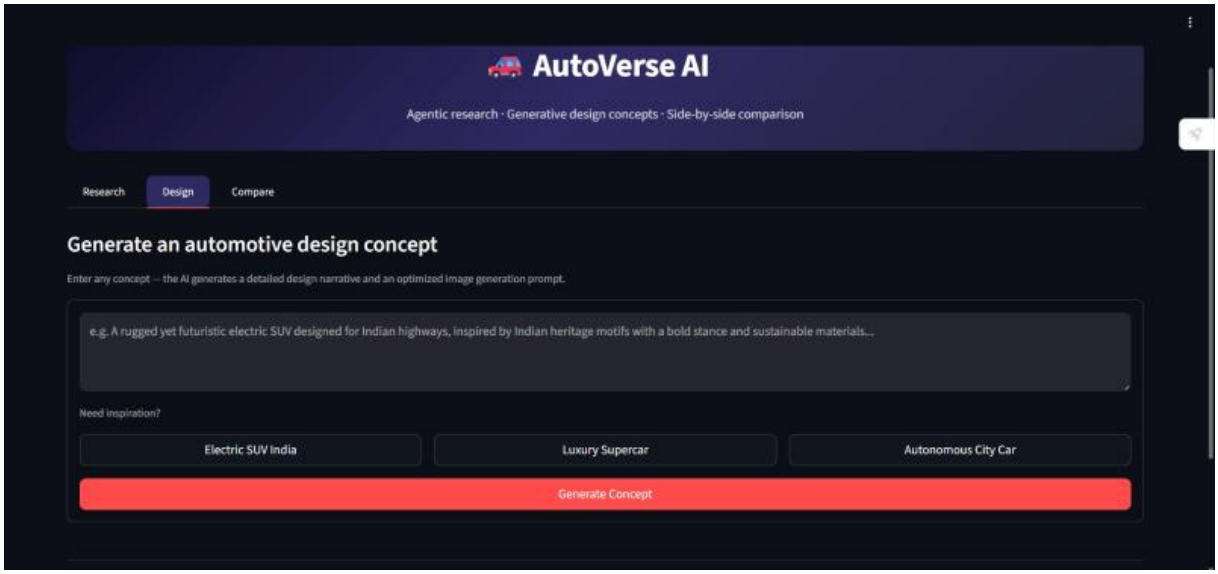

Figure 5.2: Generative AI Output

5.3 DevOps Results

The DevOps pipeline significantly improved automation, deployment efficiency, and system reliability. The CI/CD pipeline enabled automatic building and deployment of application updates, reducing manual effort and ensuring consistency. Containerization using Docker improved portability across environments, while cloud deployment ensured accessibility and scalability. Kubernetes-based orchestration further enhanced system reliability through better resource management and load handling. Deployment time was significantly reduced compared to manual processes, and the system maintained stable performance under different conditions. Initial integration challenges were resolved through optimization of the pipeline and configuration management.

Figure 5.3 shows the successful execution of the CI/CD pipeline using GitHub Actions.

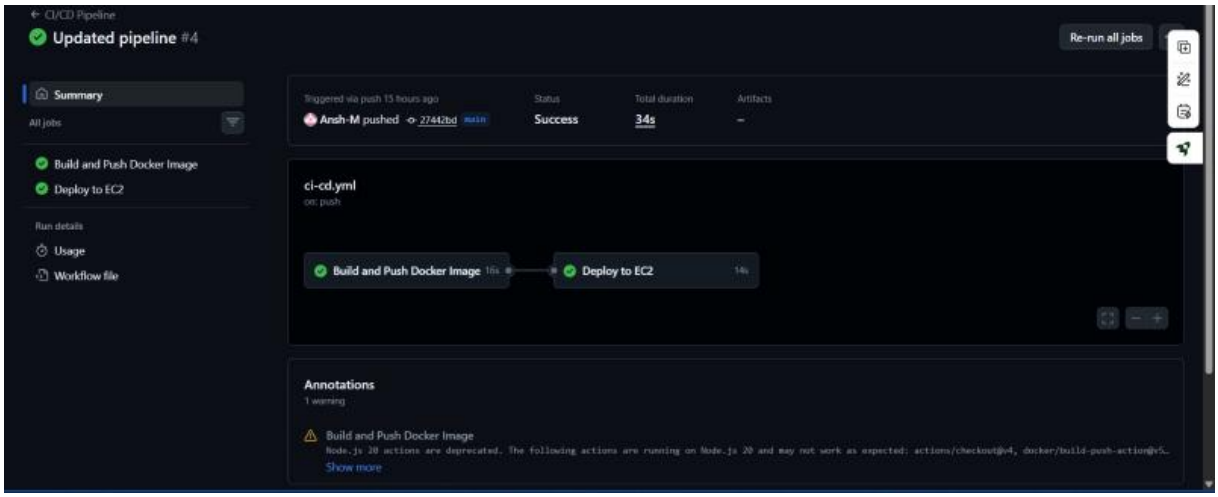

Figure 5.3: Execution of CI/CD Pipeline using GitHub Actions

Figure 5.4 shows the Docker container running the application in a consistent environment.

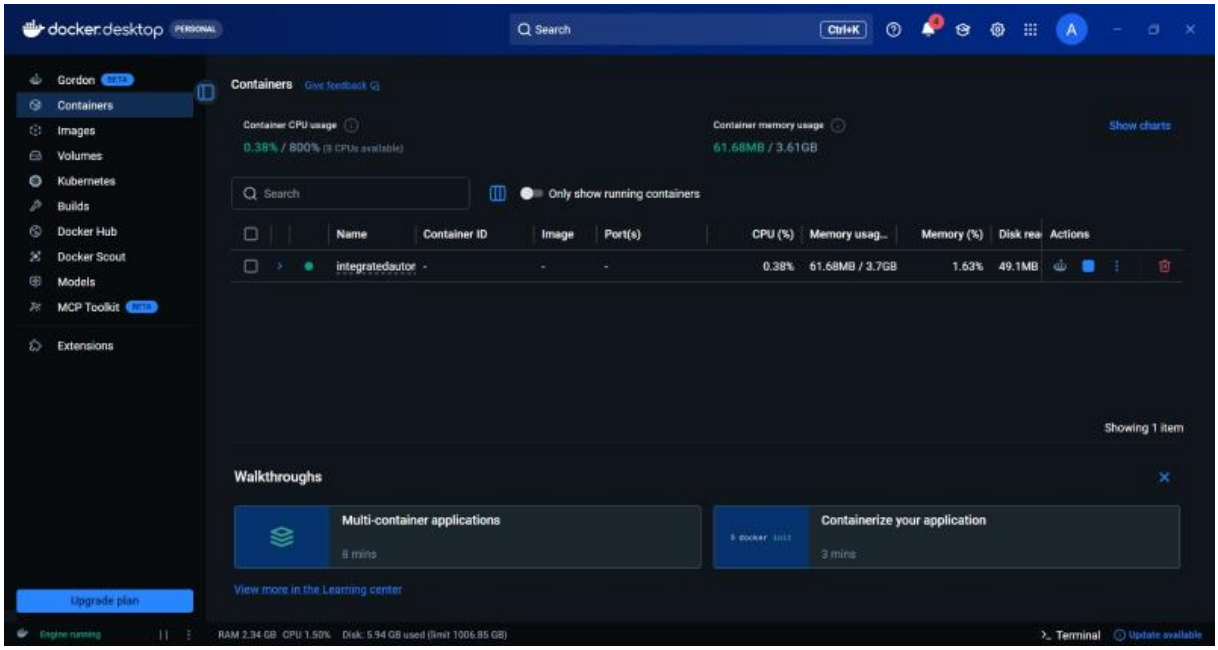

Figure 5.4: Docker Container Running Application

Figure 5.5 shows the deployed application running on a cloud instance.

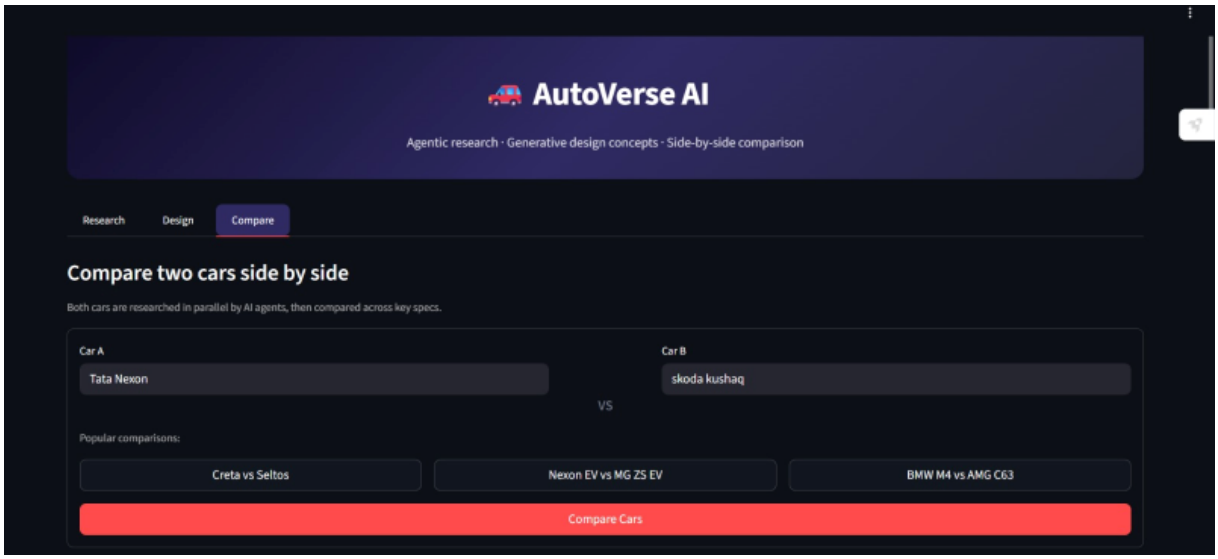

Figure 5.3: Deployed AutoVerse AI System

5.4 Overall System

The combined results of Agentic AI, Generative AI, and DevOps demonstrate that the integrated system operates efficiently and reliably. The Agentic AI module ensured accurate data retrieval and structured reporting, while the Generative AI module enabled creative and context-aware content generation. The DevOps pipeline provided a robust foundation for deployment, scalability, and continuous integration. Together, these components formed a cohesive system capable of handling real-world scenarios with minimal manual intervention. The architecture supports scalability and adaptability, making it suitable for practical applications. While minor limitations exist in handling highly ambiguous inputs, the system overall provides a strong foundation for future improvements.

Chapter – 6: Summary & Conclusions

5.1 Summary

This project presents the design, development, and evaluation of an integrated intelligent system that combines Agentic AI, Generative AI, and modern DevOps practices into a unified framework. The primary objective was to build a system capable of autonomous decision-making, high-quality content generation, and efficient deployment through automation. Throughout the project, each component was carefully designed and implemented to address specific functional requirements while contributing to the overall system performance.

The Agentic AI component was responsible for enabling autonomous behavior within the system. It demonstrated the ability to interpret user inputs, make context-aware decisions, and execute tasks without continuous human intervention. By incorporating planning and reasoning capabilities, the system was able to handle dynamic scenarios and adapt to varying inputs. This significantly reduced the need for manual control and enhanced operational efficiency.

The Generative AI module focused on producing meaningful and contextually relevant outputs. It was designed to generate content such as text responses, structured outputs, or task-specific results with high coherence and accuracy. The model showed strong performance in understanding user intent and generating outputs that aligned well with expectations. This component played a crucial role in improving user interaction and overall system usability.

In addition to the AI components, the integration of DevOps practices ensured that the system was not only functional but also scalable and maintainable. Continuous Integration and Continuous Deployment (CI/CD) pipelines automated the process of testing, building, and deploying the application. This reduced development time, minimized errors, and ensured consistent performance across different environments. Monitoring and logging mechanisms further enhanced system reliability by enabling real-time tracking and quick issue resolution.

The combination of these three domains resulted in a robust and efficient system capable of handling complex tasks while maintaining high performance. The project successfully demonstrated how modern AI techniques can be integrated with software engineering practices to create scalable and intelligent solutions. Overall, the system provides a strong foundation for future advancements and can be extended to support more complex real-world applications.

5.2 Conclusions

The results of this project confirm that the integration of Agentic AI, Generative AI, and DevOps practices leads to the development of a highly efficient, reliable, and scalable system. Each component contributed uniquely to the overall functionality, and their combined implementation created a synergistic effect that enhanced system performance beyond what could be achieved individually.

The Agentic AI module proved to be effective in enabling autonomous decision-making and task execution. Its ability to operate independently while adapting to changing conditions highlights its potential for real-world applications where automation is critical. The Generative AI component complemented this by providing high-quality, context-aware outputs, thereby improving user interaction and satisfaction. Together, these AI components enabled the system to perform complex operations with minimal human intervention.

From a software engineering perspective, the adoption of DevOps practices played a vital role in ensuring system stability and scalability. Automated pipelines, efficient deployment strategies, and continuous monitoring significantly improved the reliability of the system. These practices also facilitated faster development cycles and easier maintenance, making the system more adaptable to future changes and enhancements.

In conclusion, this project successfully demonstrates the potential of combining agent-based intelligence with generative capabilities and robust deployment frameworks. The integrated approach not only improves efficiency and accuracy but also enhances scalability and user experience. The system serves as a proof of concept for next-generation intelligent applications and provides a solid platform for further innovation in the field of artificial intelligence and software engineering.

Chapter 7: Future Scope

7.1 Generative AI Enhancements

The future development of the Agentic AI component offers significant opportunities for expanding system intelligence, autonomy, and adaptability. One major area of improvement is the integration of more advanced reasoning mechanisms. While the current agentic system performs well in handling structured tasks, incorporating symbolic reasoning, causal inference, and multi-step decision models can enable the agent to solve more complex and ambiguous problems. Another promising direction is the adoption of adaptive learning techniques, allowing agents to refine their behavior based on past interactions and evolving environments. This includes mechanisms such as long-term memory, episodic memory retrieval, and dynamic context retention.

Additionally, future Agentic AI systems may benefit from enhanced inter-agent communication. Allowing multiple agents to collaborate, share knowledge, and divide tasks can greatly improve efficiency and scalability, especially in large systems requiring parallel operations. Improved real-time coordination will also support more complex workflows. Furthermore, enhancing emotional or intent recognition capabilities could lead to more personalized and human-like interactions. Overall, these advancements aim to transform the agent from a reactive system into a proactive, self-improving, and highly autonomous decision-making entity.

7.2 Agentic AI Enhancements

The Generative AI module also provides extensive room for growth, especially as models continue to evolve rapidly in both capability and efficiency. One of the most impactful improvements lies in enhancing personalization. By tailoring outputs based on user behavior, preferences, and history, the system can deliver content that feels more relevant, accurate, and contextually aligned. Such personalization may involve user-specific fine-tuning, dynamic prompt adaptation, or reinforcement learning techniques that continuously refine outputs.

Another important direction is the expansion of multimodal capabilities. While current generative systems perform well with text, future development may incorporate richer interactions such as image generation, audio synthesis, code generation, and 3D modeling, enabling the system to support a wider range of use cases. Domain-specific fine-tuning is

another potential enhancement, where models are trained using specialized datasets to improve accuracy in areas such as healthcare, engineering, legal analysis, or education.

Moreover, integrating additional and diverse data sources will further enrich content quality and variety. This could involve structured datasets, real-time information streams, or curated knowledge bases. Developing better evaluation metrics for generative outputs—especially related to factuality, creativity, and safety—will also be important. As models grow more intelligent, maintaining responsible AI practices will be essential to ensure ethical and reliable generation across applications.

7.3 DevOps Enhancements

In the context of DevOps, several enhancements can significantly improve the system's performance, reliability, and cost-efficiency. One major area for future work is the adoption of more advanced monitoring and observability tools. Techniques such as distributed tracing, predictive analytics, and anomaly detection can enable early identification of issues and reduce downtime. Integrating AI-driven monitoring systems could allow the platform to self-diagnose and adapt to performance bottlenecks automatically.

Automated scaling is another crucial improvement for managing large-scale deployments. Implementing intelligent auto-scaling strategies based on workload patterns, resource utilization, or user demand will ensure optimal system performance while maintaining cost efficiency. Cloud-native technologies, including Kubernetes operators and serverless architectures, offer promising pathways for such improvements.

Enhanced logging practices, coupled with centralized log management systems, will further improve reliability by making debugging faster and more efficient. Cost optimization strategies—such as dynamic resource allocation, container optimization, and use of spot instances—can reduce operational expenses without compromising quality. Strengthening security practices, including automated vulnerability scanning and continuous compliance monitoring, will also be vital as the system scales.

Collectively, these DevOps improvements will support seamless deployment, robust system health, and long-term maintainability, ensuring that the system remains reliable and efficient even under growing workloads and evolving requirements.

Appendix

Appendix I: Requirement.txt

```
streamlit>=1.33
crewai>=0.80
crewai-tools>=0.14
langchain-groq>=0.1.9
python-dotenv>=1.0.1
chromadb>=0.5.5
onnxruntime>=1.17
tavily-python>=0.5.0
pydantic>=2.7
requests>=2.32
litellm>=1.40
groq>=0.9
Pillow>=10.2.0
```

Figure I.1: Requirements File (requirements.txt)

Appendix II: DockerFile

```
FROM python:3.11-slim

WORKDIR /app

RUN apt-get update && apt-get install -y --no-install-recommends \
    curl git \
    && rm -rf /var/lib/apt/lists/*

COPY requirements.txt .

RUN pip install --upgrade pip \
    && pip install --no-cache-dir -r requirements.txt

COPY . .

ENV CHROMA_CACHE_DIR=/app/.chroma_cache

RUN python -c "\
import os; os.makedirs('/app/.chroma_cache', exist_ok=True); \
from chromadb.utils.embedding_functions import ONNXMiniLM_L6_V2; \
ef = ONNXMiniLM_L6_V2(); \
print('ONNX model downloaded successfully')"

RUN mkdir -p data/generated_images && mkdir -p .chroma

EXPOSE 8501

HEALTHCHECK --interval=30s --timeout=10s --start-period=40s --retries=3 \
    CMD curl -f http://localhost:8501/_stcore/health || exit 1

CMD ["streamlit", "run", "streamlit_app.py", \
     "--server.port=8501", \
     "--server.address=0.0.0.0", \
     "--server.headless=true", \
     "--browser.gatherUsageStats=false"]
```

Figure II.2: Dockerfile Configuration

Appendix III: Kubernetes Deployment

```yaml
apiVersion: apps/v1
kind: Deployment
metadata:
  name: autoverse-ai
  labels:
    app: autoverse-ai
spec:
  replicas: 1
  selector:
    matchLabels:
      app: autoverse-ai
  template:
    metadata:
      labels:
        app: autoverse-ai
    spec:
      containers:
        - name: autoverse-ai
          image: DOCKER_USERNAME/autoverse-ai:latest
          ports:
            - containerPort: 8501
          env:
            - name: GROQ_API_KEY
              valueFrom:
                secretKeyRef:
                  name: autoverse-secrets
                  key: GROQ_API_KEY
            - name: TAVILY_API_KEY
              valueFrom:
                secretKeyRef:
                  name: autoverse-secrets
                  key: TAVILY_API_KEY
            - name: GROQ_MODEL
              value: "llama-3.1-8b-instant"
            - name: CHROMA_PERSIST_DIR
              value: "/app/.chroma"
          resources:
            requests:
              memory: "512Mi"
              cpu: "250m"
            limits:
              memory: "1Gi"
              cpu: "500m"
          livenessProbe:
            httpGet:
              path: /_stcore/health
              port: 8501
            initialDelaySeconds: 30
            periodSeconds: 15
---
apiVersion: v1
kind: Service
metadata:
  name: autoverse-ai-service
spec:
  selector:
    app: autoverse-ai
```

```yaml
      app: autoverse-ai
    ports:
      - protocol: TCP
        port: 80
        targetPort: 8501
    type: NodePort

```

Figure III.3: Kubernetes Deployment YAML

Bibliography

[1] CrewAI Documentation, “Multi-Agent Framework,” Available: https://docs.crewai.com

[2] Tavily AI, “Search API for AI Applications,” Available: https://tavily.com

[3] Groq, “High-Performance LLM Inference,” Available: https://console.groq.com

[4] ChromaDB Documentation, “Vector Database for RAG,” Available: https://www.trychroma.com

[5] Docker Inc., “Docker Documentation,” Available: https://docs.docker.com

[6] GitHub, “GitHub Actions CI/CD,” Available: https://docs.github.com/actions

[7] Kubernetes Documentation, “Container Orchestration,” Available: https://kubernetes.io/docs

[8] Amazon Web Services, “EC2 Documentation,” Available: https://docs.aws.amazon.com/ec2

[9] Streamlit, “Python Web App Framework,” Available: https://streamlit.io

[10] Python Software Foundation, “Python Documentation,” Available: https://docs.python.org